

# 01\_player\_consistency

December 18, 2025

## 1 Q1: Player Consistency in Scoring, Rebounding, and Assists

The goal of this notebook is to identify the most consistent players across key performance metrics.

We define consistency as low game-to-game variability: - We compute the mean and standard deviation of PTS, TRB, and AST across all their games. - Players with the lower standard deviation, given a reasonable number of games, are considered more consistent.

I plan to: 1. Load the game-level data 2. Filter by players with a minimum number of games played 3. Compute per-player mean and std of PTS, TRB, and AST 4. Rank players by consistency 5. Visualize the most consistent players using boxplots and scatterplots

Imports for future plotting and data manipulation.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from utils import add_gamekey_and_win, add_per_minute_stats

df = pd.read_csv("data/player_game_stats_clean.csv")
```

Filter the dataset to only include players with enough games. I chose 30 games because it seems like a reasonable amount of games to get a good variability of statistics from. We don't want players that have played only 1 game and never varied. I set a variable for min\_games so we can easily adjust it as well.

```
[2]: min_games = 30

games_per_player = df.groupby("Player")["PTS"].count().rename("GamesPlayed")
games_per_player.head()

# Merge games played into the main df
df = df.merge(games_per_player, on="Player")

# Filter
df_filtered = df[df["GamesPlayed"] >= min_games].copy()
print("Original players:", games_per_player.shape[0])
```

```
print("Players with at least", min_games, "games:", df_filtered["Player"] .
    ↪nunique())
```

Original players: 562  
 Players with at least 30 games: 305

Now I compute the consistency metrics per player by grouping by player and team to get their PTS, TRB, and AST mean and standard deviation.

```
[3]: player_consistency = (
    df_filtered
    .groupby(["Player", "Tm"])
    .agg(
        GamesPlayed=("PTS", "count"),
        PTS_mean=("PTS", "mean"),
        PTS_std=("PTS", "std"),
        TRB_mean=("TRB", "mean"),
        TRB_std=("TRB", "std"),
        AST_mean=("AST", "mean"),
        AST_std=("AST", "std")
    )
    .reset_index()
)

player_consistency["PTS_cv"] = player_consistency["PTS_std"] / □
    ↪player_consistency["PTS_mean"].replace(0, np.nan)
player_consistency["TRB_cv"] = player_consistency["TRB_std"] / □
    ↪player_consistency["TRB_mean"].replace(0, np.nan)
player_consistency["AST_cv"] = player_consistency["AST_std"] / □
    ↪player_consistency["AST_mean"].replace(0, np.nan)

player_consistency.head()
```

	Player	Tm	GamesPlayed	PTS_mean	PTS_std	TRB_mean	TRB_std	AST_mean	AST_std	PTS_cv	TRB_cv	AST_cv
0	A.J. Green	MIL	44	7.659091	5.382760	2.250000	1.780057	1.272727	1.318273	0.702794	0.791137	1.035786
1	Aaron Gordon	DEN	30	12.333333	6.608946	4.733333	2.981938	3.066667	2.531639	0.535860	0.629987	0.825534
2	Aaron Holiday	HOU	36	4.222222	3.742506	0.944444	1.067559	1.194444	1.214659	0.886383	1.130357	1.016924
3	Aaron Wiggins	OKC	51	10.137255	6.720177	3.568627	2.492027	1.568627	1.374844	0.662919	0.698315	0.876463
4	Adem Bona	PHI	36	3.000000	3.144156	2.833333	2.489980	0.305556	0.524783	1.048052	0.878816	1.717470

Now I rank players by their consistency, which I define as the smallest standard deviation in each

statistic. I .head(10) to limit the top 10 most consistent scorers from PTS, rebounders from TRB, and playmakers from AST

```
[4]: min_games = 30

player_consistency_filtered = player_consistency[
    player_consistency["GamesPlayed"] >= min_games].copy()

# Top 10 most consistent scorers
top_consistent_pts = (
    player_consistency_filtered
    .sort_values(["PTS_std", "PTS_mean"], ascending=[True, False])
    .head(10))

# Top 10 most consistent rebounders
top_consistent_trb = (
    player_consistency_filtered
    .sort_values(["TRB_std", "TRB_mean"], ascending=[True, False])
    .head(10))

# Top 10 most consistent playmakers/assists
min_ast_mean = 2.0

pc = player_consistency[player_consistency["GamesPlayed"] >= min_games].copy()
ast_cutoff = pc["AST_mean"].quantile(0.75)
playmakers = pc[pc["AST_mean"] >= ast_cutoff].copy()

top_consistent_ast = (
    playmakers
    .sort_values(["AST_std", "AST_mean"], ascending=[True, False])
    .head(10)
)

top_consistent_ast[["Player", "Tm", "GamesPlayed", "AST_mean", "AST_std"]]
```

```
[4]:      Player    Tm  GamesPlayed   AST_mean   AST_std
116  Isaiah Hartenstein  OKC          31  4.129032  1.802627
148       Jaylen Brown   BOS          45  4.777778  1.832644
37        CJ McCollum   NOP          38  3.815789  1.872349
45        Caris LeVert   CLE          38  3.684211  1.918689
99       Dyson Daniels  ATL          47  3.957447  1.921927
89     Donovan Mitchell   CLE          49  4.734694  1.955500
135       Jalen Suggs   ORL          35  3.685714  2.011313
136       Jalen Williams  OKC          48  5.083333  2.019409
201       Kyrie Irving   DAL          41  4.804878  2.039847
287  T.J. McConnell   IND          49  4.510204  2.072865
```

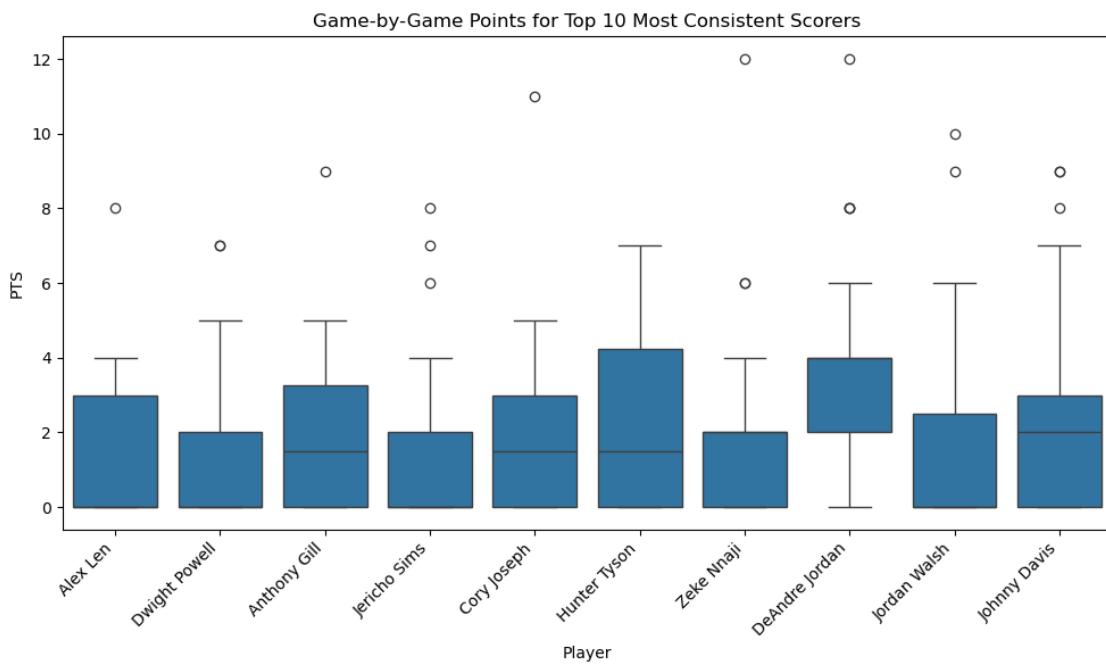
Using seaborn to visualize game-by-game distributions for the top 10 most consistent scorers.

```
[10]: # Get the names of the top 10 consistent scorers
top_pts_players = top_consistent_pts["Player"].tolist()

df_top_pts = df_filtered[df_filtered["Player"].isin(top_pts_players)].copy()

plt.figure(figsize=(10, 6))
sns.boxplot(
    data=df_top_pts,
    x="Player",
    y="PTS",
    order=top_pts_players)

plt.xticks(rotation=45, ha="right")
plt.title("Game-by-Game Points for Top 10 Most Consistent Scorers")
plt.ylabel("PTS")
plt.xlabel("Player")
plt.tight_layout()
plt.show()
```



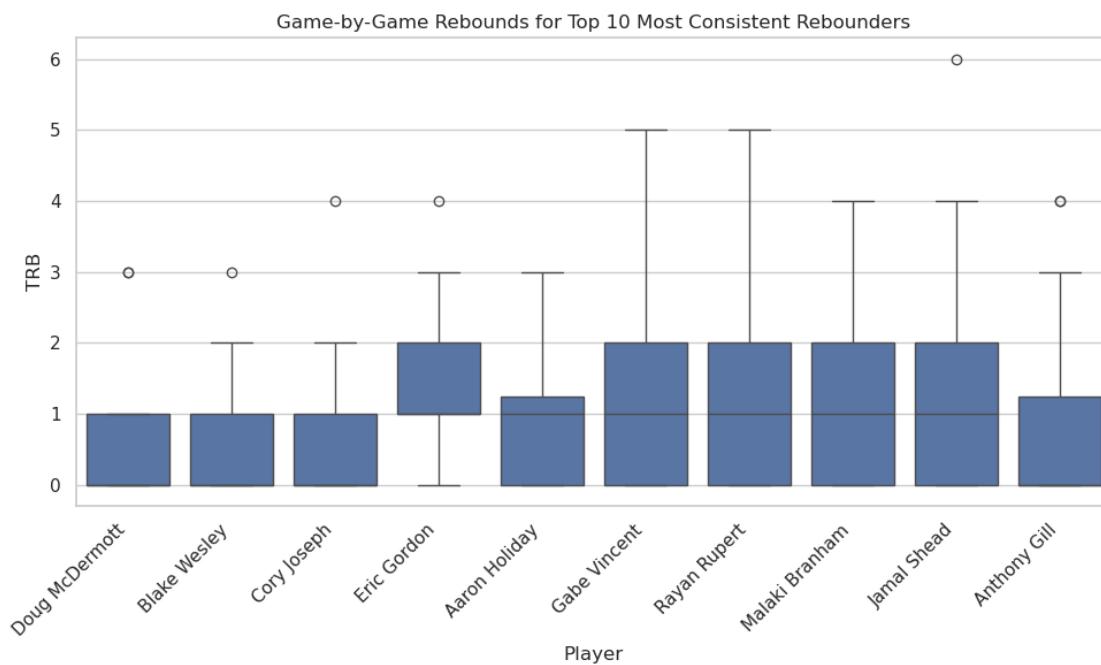
## 2

Using seaborn to visualize game-by-game distributions for the top 10 most consistent rebounders and playmakers, similar to the plot above.

```
[34]: # Top 10 rebounders by consistency
top_trb_players = top_consistent_trb["Player"].tolist()
df_top_trb = df_filtered[df_filtered["Player"].isin(top_trb_players)].copy()

plt.figure(figsize=(10, 6))
sns.boxplot(
    data=df_top_trb,
    x="Player",
    y="TRB",
    order=top_trb_players)

plt.xticks(rotation=45, ha="right")
plt.title("Game-by-Game Rebounds for Top 10 Most Consistent Rebounders")
plt.ylabel("TRB")
plt.xlabel("Player")
plt.tight_layout()
plt.show()
```



```
[31]: top_ast_players = top_consistent_ast["Player"].tolist()

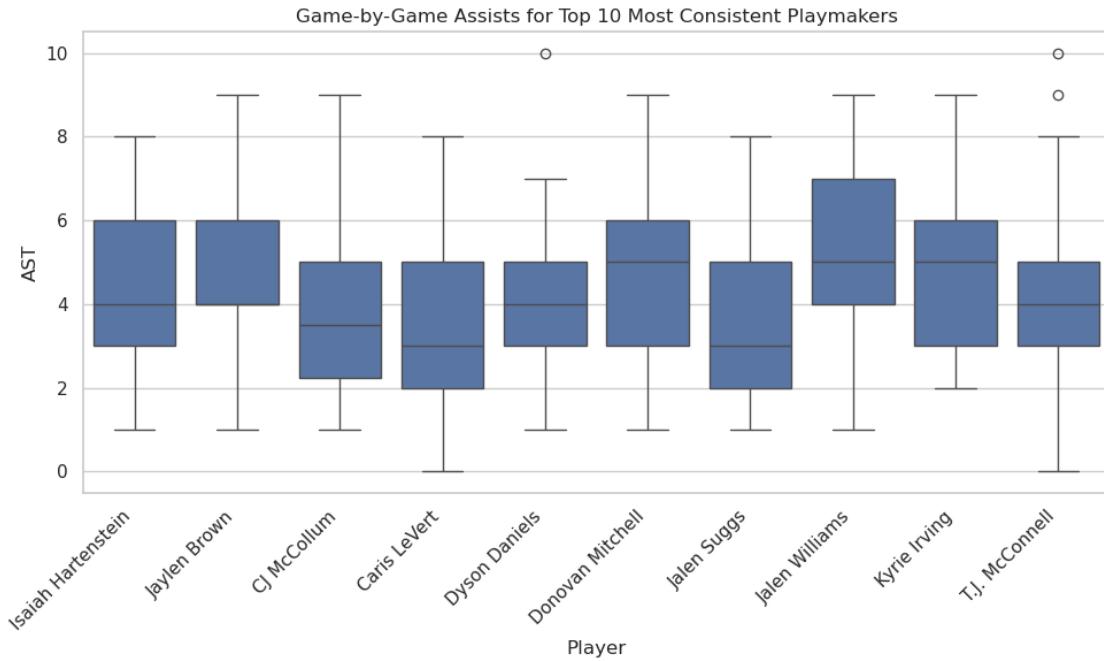
df_top_ast = df_filtered[df_filtered["Player"].isin(top_ast_players)].copy()

plt.figure(figsize=(10,6))
sns.boxplot(data=df_top_ast, x="Player", y="AST", order=top_ast_players)
plt.xticks(rotation=45, ha="right")
```

```

plt.title("Game-by-Game Assists for Top 10 Most Consistent Playmakers")
plt.tight_layout()
plt.show()

```



Mean versus Variability scatterplots to visualize the tradeoff between volume and consistency for points, rebounds, and assists.

```

[32]: plt.figure(figsize=(7, 5))
sns.scatterplot(
    data=player_consistency,
    x="PTS_mean",
    y="PTS_std",
    alpha=0.6
)
plt.title("Points: Mean vs Standard Deviation (Player-Level)")
plt.xlabel("Average Points per Game")
plt.ylabel("Standard Deviation of Points")
plt.tight_layout()
plt.show()

plt.figure(figsize=(7, 5))
sns.scatterplot(
    data=player_consistency,
    x="TRB_mean",
    y="TRB_std",

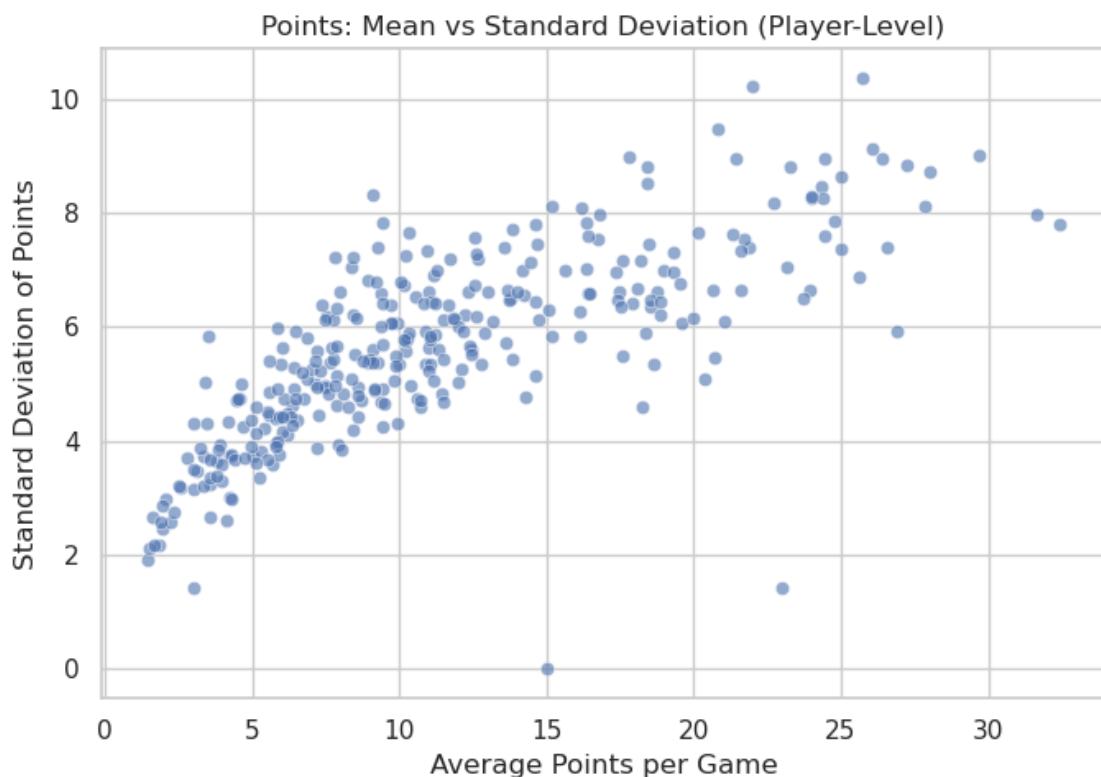
```

```

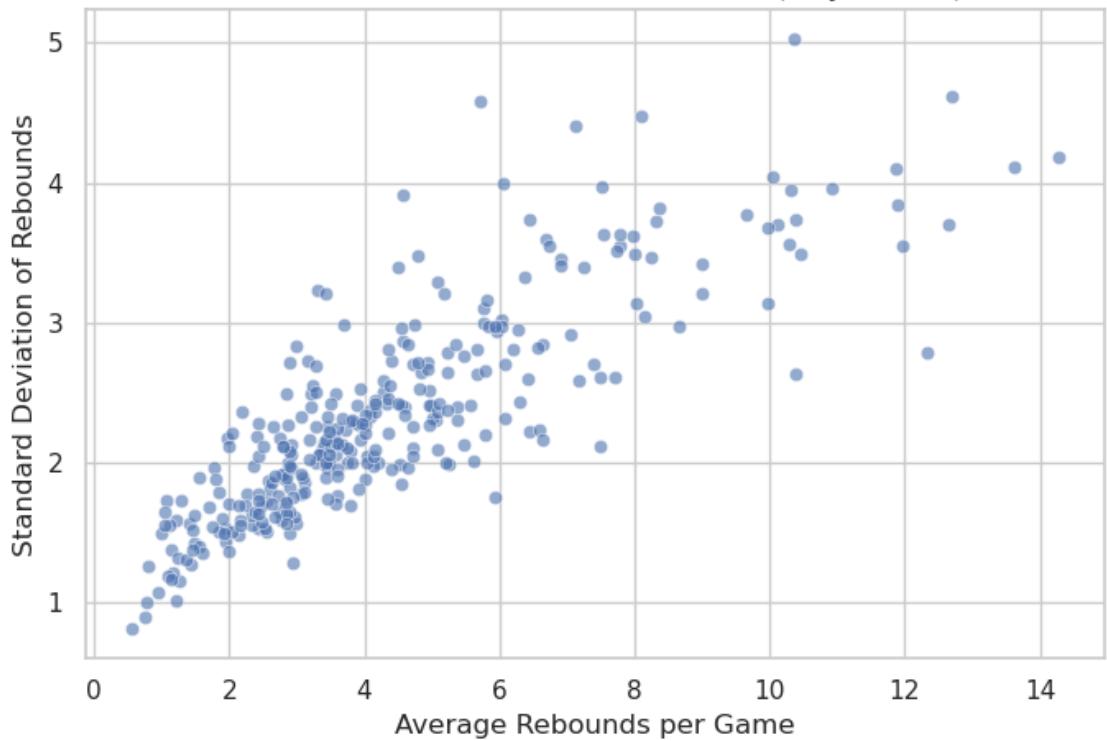
        alpha=0.6
    )
plt.title("Rebounds: Mean vs Standard Deviation (Player-Level)")
plt.xlabel("Average Rebounds per Game")
plt.ylabel("Standard Deviation of Rebounds")
plt.tight_layout()
plt.show()

plt.figure(figsize=(7, 5))
sns.scatterplot(
    data=player_consistency,
    x="AST_mean",
    y="AST_std",
    alpha=0.6
)
plt.title("Assists: Mean vs Standard Deviation (Player-Level)")
plt.xlabel("Average Assists per Game")
plt.ylabel("Standard Deviation of Assists")
plt.tight_layout()
plt.show()

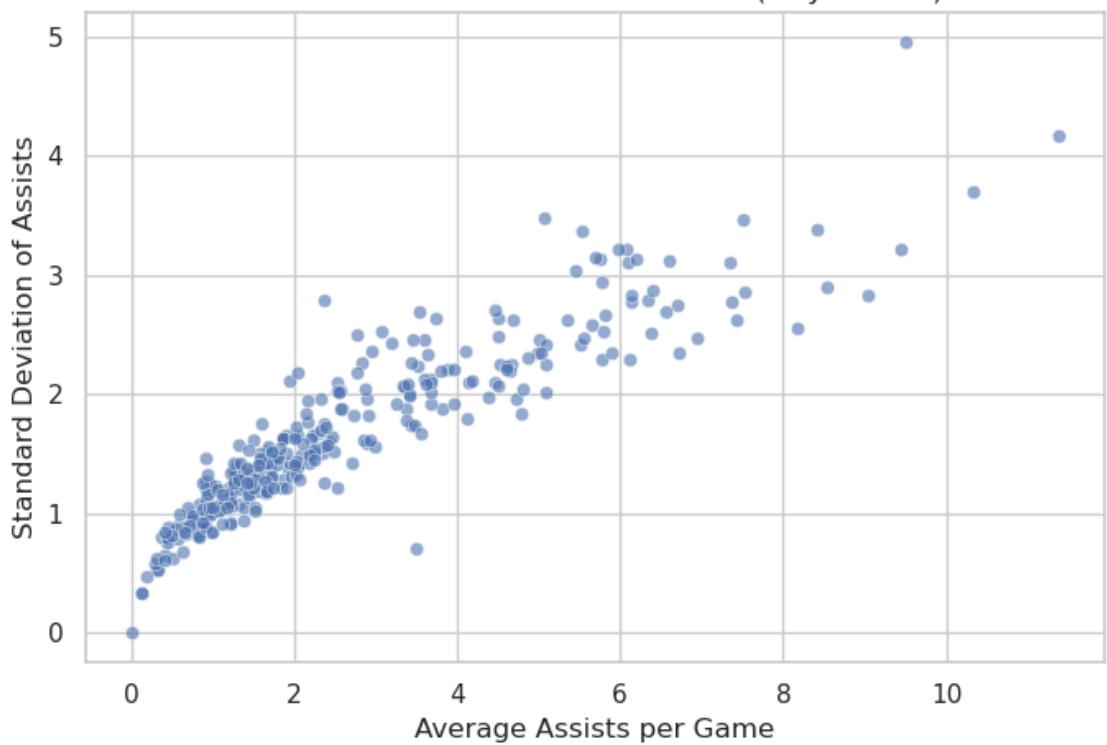
```



Rebounds: Mean vs Standard Deviation (Player-Level)



Assists: Mean vs Standard Deviation (Player-Level)



[ ]: